

More Advanced GA Concepts

Seth Bullock

bristol.ac.uk

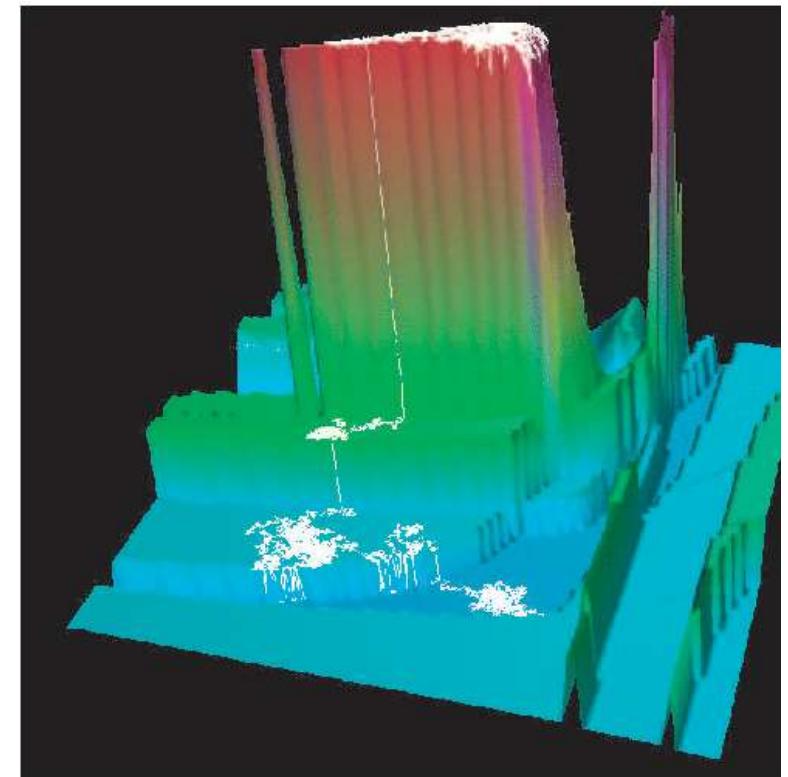
This lecture covers a number of Evolutionary Computation concepts:

- Premature Convergence
- Neutrality
- Quasi-Species
- Epistasis
- Problem Modularity
- No Free Lunch
- Evolvability

The Problem of Local Optima

- Hill-climbers get stuck on local optima; Exhaustive search does not
- GAs are less likely to get stuck on local optima than hill-climbers:
 - ...because they maintain and evolve a *population of solutions*
 - ...because genetic operators may generate a *wider range of neighbours*
 - These reasons exploit and rely on *diversity* in the evolving population
- Local optima are still a problem as *pop diversity is not guaranteed*
- Should a population of solutions end up *converged* within one **basin** of attraction they can find it difficult to escape
 - Getting stuck like this is called: *premature convergence*

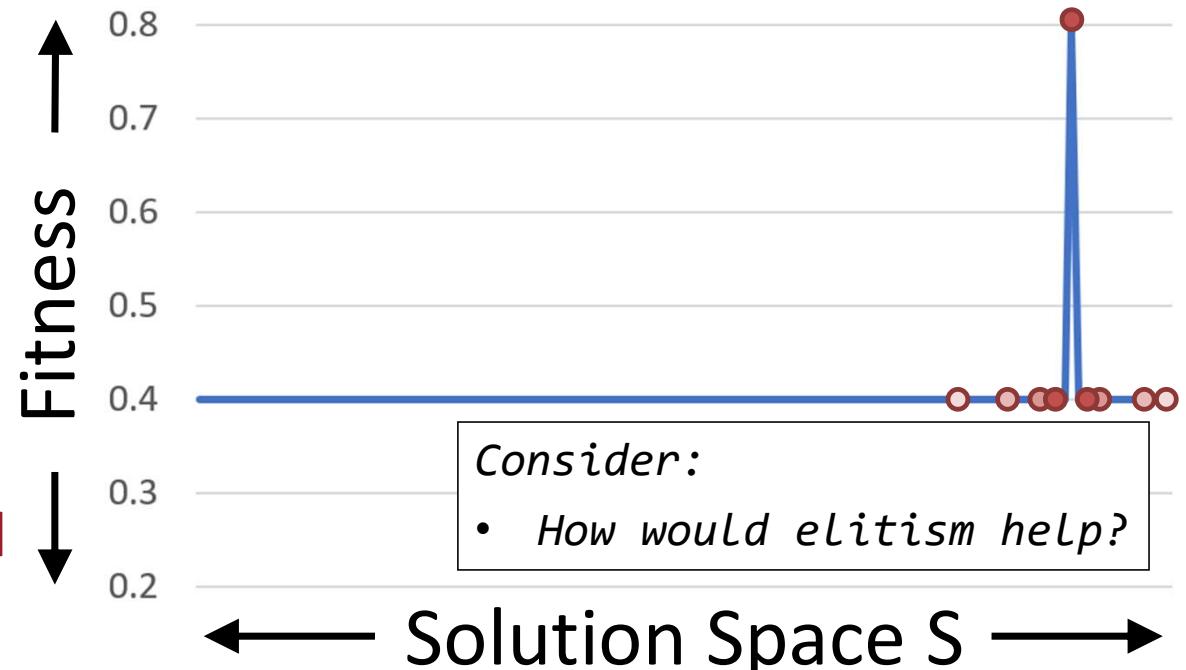
- Neighbouring genotypes with equal fitness are ‘selectively neutral’
 - ...even if their phenotypes are different.
 - Evolution cannot choose between them...
 - ...which results in ‘evolutionary drift’
- Neutrality is often easy to overlook and hard to visualize, but it can be key...
 - E.g., ‘neutral networks’ that percolate the search space may allow converged populations to escape “local optima”



Barnett (2002).
[Explorations in Evolutionary Visualisation](#)

The Invisible Needle

- Recall the “needle” fitness landscape we saw in a previous lecture
- The needle is hard to find because there’s no gradient to climb
- But even if we start the pop *on the needle*, we still might not find it...
 - Notice: it’s very rare for an offspring to be an exact copy of their parent
 - So a perfect solution will tend to have unfit offspring...



The Quasi-Species Concept

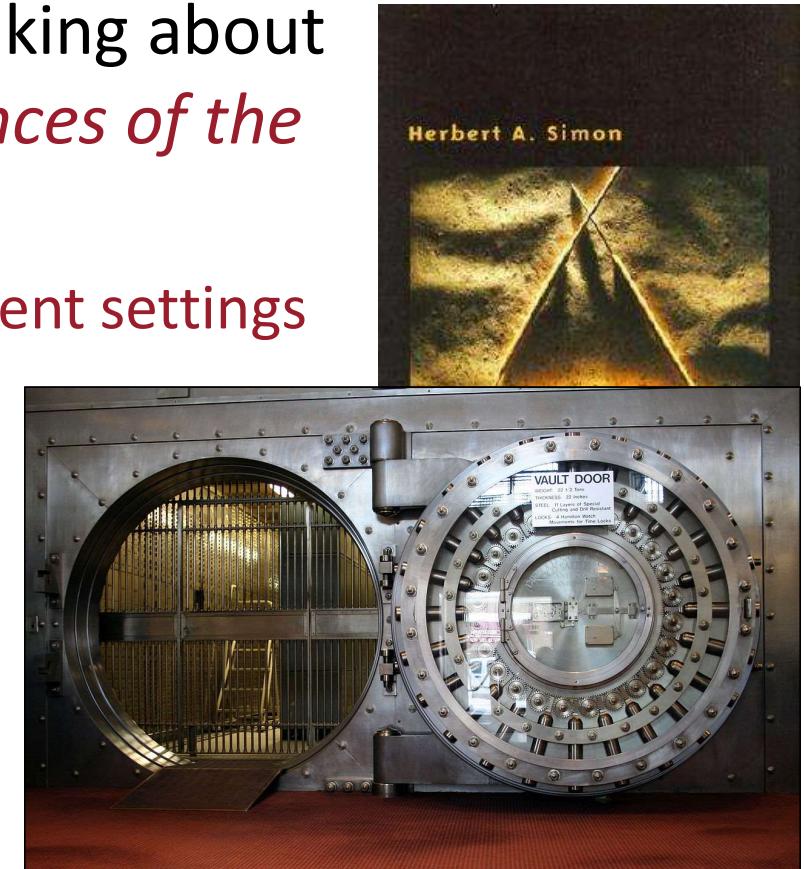
- Manfred Eigen & Peter Schuster developed the *quasi-species* concept to explain what's happening here:
 - High GA mutation rates (much higher than for DNA) mean that the evolving population is a *cloud* of points on the landscape: a *quasi-species*
 - If the selection pressure on the population to reward good solutions...
 - ...is overcome by the mutation pressure that constantly corrupts them...
 - ...then the population effectively cannot even “see” the needle.
- Conversely: if an evolving population *is* able to find a good solution, *and stay there*, we can infer that the solution must be surrounded by other pretty good solutions – it is *robust*

- If we're very lucky, fitness is a *linear function* of the gene alleles:
 - The contribution of each gene to fitness is *independent* of other genes
 - 1-max is a linear fitness function: fitness = the number of 1s in a bitstring
 - Trivial to optimize: optimize each gene independently
- But almost always, fitness is a *non-linear function* of gene alleles:
 - The contribution of genes to fitness is *interdependent* to some extent
 - The best allele choice at one gene *depends* on the alleles at other genes
 - E.g., baby skull size \leftrightarrow size of birth canal; weapon range \leftrightarrow visual range
 - This is *epistasis*: more epistasis makes a search problem harder

- The presence of epistasis \Leftrightarrow the existence of local optima.
 - Epistasis / local optima are different ways of describing the same thing:
 - Non-linearities in the mapping of fitness onto the search space
- Epistasis means we can't optimize each gene *independently*
 - Instead our algorithm must *co-optimize* the alleles of multiple genes
- How big is the set of interacting genes? (the 'order' of interaction)
- How many sets of interacting genes? How do they overlap?
- More generally, what is the *structural modularity* of the problem?

Modular Decomposability

- Herb Simon introduces a nice way of thinking about problem modularity in his book *The Sciences of the Artificial* (1969): trying to crack a safe
 - A bank safe has N dials, each with 100 different settings
 - Only *one* setting is correct on each dial
 - *All* dials need to be correct to open the safe
 - (Recall that Dawkins uses the same example in his *Horizon* episode...)
 - How hard is it to open the safe?



Modular Decomposability

The answer depends on the *modularity* of the safe

- For a safe with perfectly silent dials: we must search all 100^N possibilities
 - We can expect to have to check $\frac{1}{2}$ of them => $100^N/2$ to crack the safe
 - No modularity; no useful problem structure; solution has strong epistasis
- But if each dial gives a little *click* only when it is at the correct setting:
 - We can expect to have to check $\frac{1}{2}$ of the 100 possibilities *for each dial*
 - => $100N/2$ tries to crack the whole safe (much less time)
 - Full modularity; useful problem structure; solution has zero epistasis
 - The same as the “Methinks it is like a weasel” & 1-max problems

Partial Decomposability

- But real problems tend to lie somewhere in between:
 - Partial modularity; useful problem structure; solution has some epistasis
 - Compare: “Methinks...” vs. evolving any correct English sentence
 - English sentences have modules (words) that depend on each other
- Watson (2006): consider a safe with dials that are sensitive to each other. For each dial, n settings (including the correct one) give a little click:
 - $n=1$ (full modularity); $n=100$ (no modularity); $n=20$ (some modularity)
 - We must try $n^N/2$ combinations (much less than $100^N/2$) to crack the safe
 - (After we find which are the n clicky settings on each dial: $100N$ tries)
 - What if n is halved for every dial that is in the correct setting?

No Free Lunches

- Wolpert and Macready (1997): No Free Lunch for Optimization:
 - When algorithm performance is averaged across *all possible problems*:
 - *Any two optimization algorithms will exhibit equivalent performance*
 - *i.e., no optimization algorithm can do better than blind random search*
 - The insight here is that every search algorithm other than random search has a *search bias*: it works on a hunch, gamble, or heuristic
 - For every time the hunch pays off, there's a time where it doesn't
 - Perhaps need to shift focus to *satisficing* rather than optimising...
 - Algorithm quality depends on the *problem structure* that it faces
-

Learning How to Search

- Recall that the structure of S (including its modularity) is influenced by our choice of *representation* and *genetic operators*
- So, we can restructure the space by making different choices
- Good choices could be the difference between success and failure
 - Might it be possible to learn which operators to use when?
 - Or to learn a good genetic representation?
 - ..either for a whole class of different, related problems (e.g., timetables)?
 - ..or *online* for one problem *as we are trying to solve it*
 - Could a GA implement evolution that gets more + more powerful?

The Evolution of Evolvability

- Natural evolution has got more powerful over time:
 - cf. the ‘major transitions’ in evolution (Maynard Smith and Szathmáry)
 - E.g., Single=>Multi-cellularity; A-sex=>Sex; Non-Social=>Social/Cultural
- In fact, Dawkins’ paper at the first Artificial Life conference is on what his *Biomorphs* tell us about “the Evolution of Evolvability”.
 - ‘The ability of a population to generate *adaptive* genetic diversity’
 - (Here *adaptive* means ‘well adapted’ to the problem at hand)
 - i.e., evolvability is: how good is a population at evolving
 - A poor GA has low evolvability; Nature exhibits high evolvability

The Evolution of Evolvability

- More subtle mechanisms in nature:
 - DNA repair – maintains some parts of the genome more carefully
 - Chromosome Organisation – nearby genes less likely to be disrupted
- Genetic Algorithms can look to exploit similar tricks:
 - Analyse how fitness varies in the current population – build a model
 - Use this model to predict what type of genotypic variation will be best
 - Effectively learning the structure of the problem during evolution
- E.g., “How can evolution learn?” (Watson & Szathmáry, 2015)
 - <http://www.bbc.com/earth/story/20170301-life-may-actually-be-getting-better-at-evolving>

Example Questions

- Complete the missing fitness entries in the two tables of genotypes and their associated fitness values, below, such that there is (i) no epistasis, (ii) some epistasis. [2 marks]
- Give an example of two aspects of a real-world problem that are linked epistatically. Explain your answer. [4 marks]
- A problem comprises two complex but almost independent sub-problems. How should a GA's genotype be structured, and what type of crossover should be used? [6 marks]

Thank you!

Coevolution

Seth Bullock

bristol.ac.uk

Fitness Function Design Is Hard

- As we have seen, a GAs success or failure is often down to our design decisions:
 - Our choice of genetic representation of the phenotype space
 - Our choice of genetic operators and selection scheme
 - Our design of a fitness function
- The third of these is particularly trying:
 - Many fitness functions are hard for GAs to make progress on due to local optima, ruggedness, neutrality, deception, etc.

- But what if we could skip that tricky bit?
- Real evolving populations *do not consult a fitness function* to find out how many offspring each individual is allowed.
 - They just get on with it
 - Organisms co-operate & compete with the other organisms in their environment, and have offspring as part of this activity
- Coevolution: *when multiple populations of evolving organisms influence each other's evolution*

- Coevolution = An automatic fitness function:
 - There is no fitness function designer in nature
- Coevolution = An implicit fitness function:
 - Instead of being set an *explicit* objective target by the fitness function, real organisms are assessed *relative* to their current environment
- Coevolution = A dynamic fitness function:
 - Instead of being assessed against a *fixed* target, real organisms cope with environments that change/evolve over time

Example

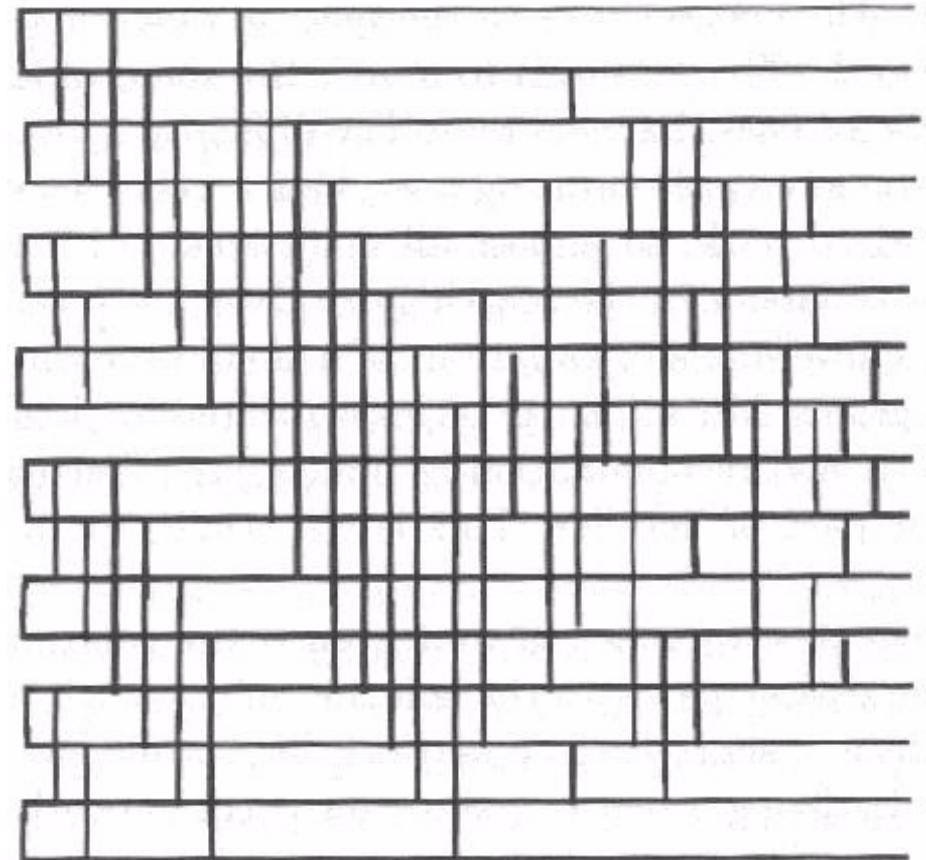
- We would like to evolve a robot goalkeeper for RoboSoccer
 - Designing a fitness function has proven to be quite hard
- Instead, we decide to *coevolve* a population of goalkeeper robots against a population of robot penalty takers
 - We start off with random penalty takers and random goalkeepers
 - As soon as penalty takers start to be able to beat random goalkeepers..
 - ..there is pressure on the goalkeepers to get better..
 - ..putting pressure on the penalty takers to get better.. ..and so on.
 - If things work out, coevolution generates a good solution without us having to define a ‘good’ strategy via an explicit fitness function

- Co-operative Coevolution:
 - Each population deals with a separate part of the overall problem
 - A solution combines individuals from each population
 - Appropriate when the overall problem is highly structured
- Competitive Coevolution:
 - One population of “solutions” to the problem that we are interested in coevolves against a population of “challengers” that must be defeated
 - Solutions are fit if they defeat many challengers.
 - Challengers are fit if they defeat many solutions.
 - Reminiscent of natural host-parasite or predator-prey coevolution.

- The first coevolutionary GA was due to Hillis (1990) and was applied to solving a list-sorting problem.
 - What's the most efficient way to sort a list of length n ?
 - Hillis pitched a population of 'hosts' (sorting networks) against a population of 'parasites' (each a set of lists to be sorted)
 - He initialised his host population with a reasonable solution and initialised his parasite population with random lists
 - Hosts rewarded for sorting parasite lists; parasites for being unsorted
 - In a few hours, he was able to coevolve a sorter that used 61 'gates'
 - Engineers had taken years to achieve a sorter with 60 gates

- History of $n=16$ sorting networks
 - Batcher & Knuth (1964) **63** gates
 - Shapiro (1969) **62** gates
 - Green (1969) **60** gates (best known)
 - Hillis (1990): Evolved **65** gates...
Hillis (1990): Coevolved **61** gates...
- Used in telecoms switches
- Now also relevant to GPUs
 - Saving even one gate can make a huge performance difference

List Sorting Example



Why Did Regular Evolution Fail?

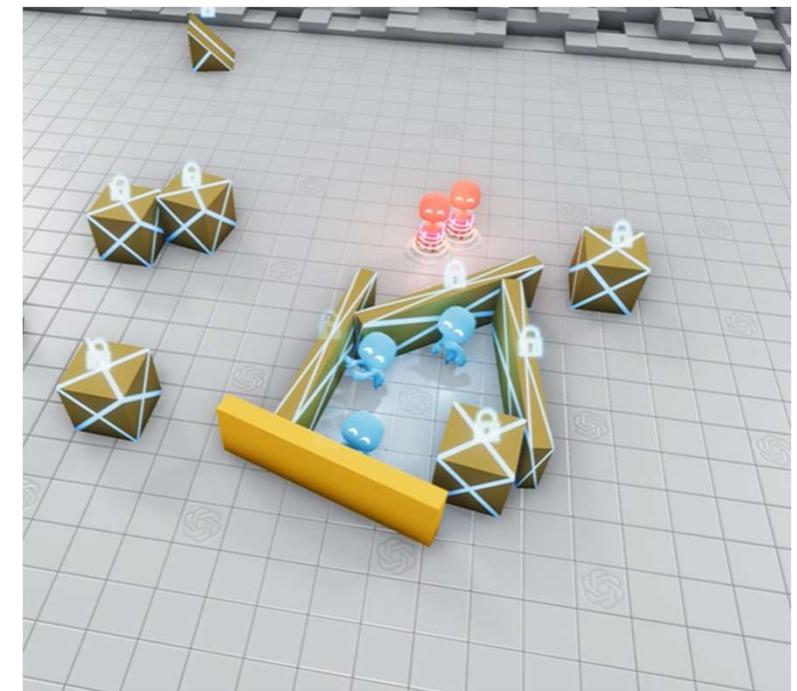
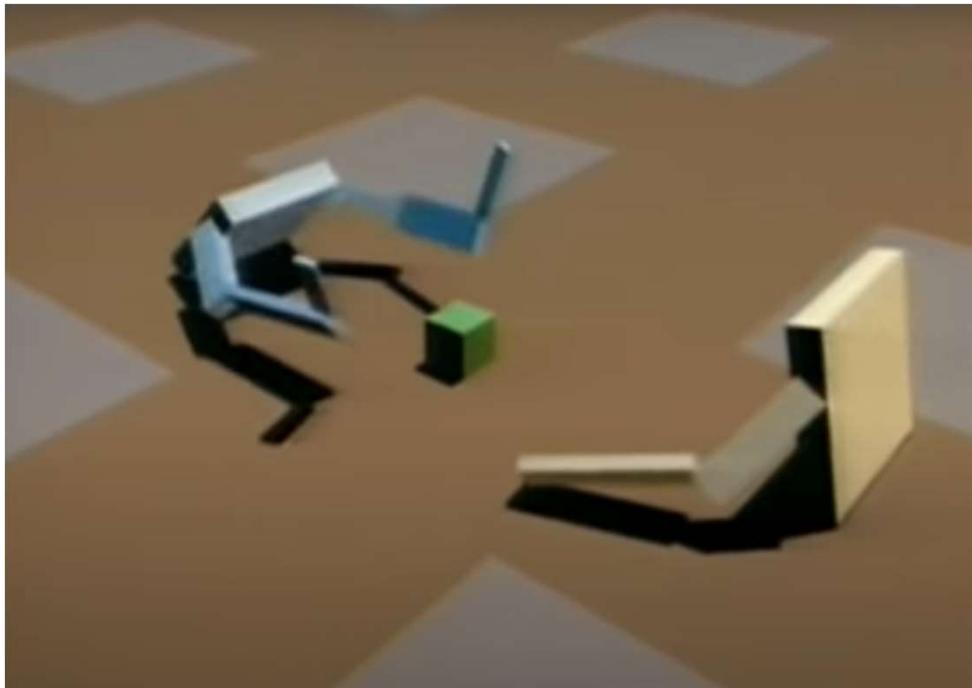
Evolution (~65K hosts for 5K generations) had problems:

1. Many local optima: reasonable sorters are surrounded by rubbish
2. Exhaustive assessment is prohibitively slow
 - There are $16!$ different lists to consider...
 - (Even reducing to 2^{16} binary lists is still a lot to check for each sorter)
3. But assessing against a random sample of lists is also inefficient:
 - Most randomly generated lists are sorted by reasonable sorters
 - Really challenging lists are rare, so there's a lack of pressure to improve

Why Was Coevolution So Much Better?

- Coevolution solved these problems automatically:
 - Instead of random lists, sorters needed to sort sets of lists that were selected to become harder and harder.
 - If sorters got stuck on a local optimum, parasites were under pressure to punish that strategy by finding the lists that it couldn't sort.
 - When improvements in the sorter population meant selection pressure got weaker, parasites were selected to get still harder.

From Karl Sims' (1994) *Blockies* to OpenAI's (2019) *Hide & Seek*:



Thank you!

Coevolution II

Seth Bullock

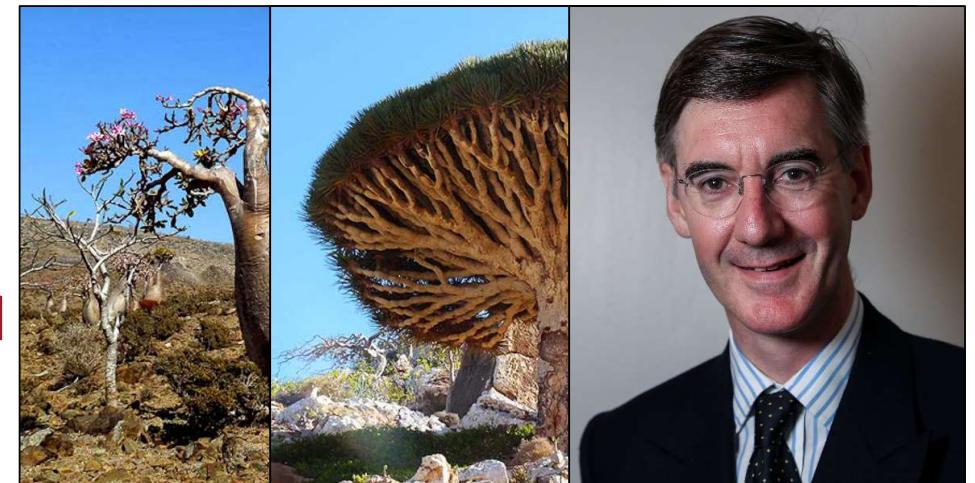
bristol.ac.uk

- Coevolutionary arms races are supposed to *escalate* to increasing levels of sophistication... but sometimes that doesn't happen:
 - Stagnation: populations reach a poised state in which there is a turnover of individuals in the population but no phenotypic progress
 - Cycling: populations repeatedly cycle through the same weak strategies as if they were playing a game of scissors-paper-stone (the Red Queen)
 - Disengagement: every host is beaten by every parasite, all selection pressure is extinguished and coevolution effectively ceases (see Lab 2)
 - Measuring Progress: without a fitness function it can be difficult to assess whether genuine progress is being made by coevolving populations.

- To a large extent the problems with coevolution are familiar:
 - Our old foe: the loss of genetic and phenotypic diversity
- How can we stop the populations from converging and becoming vulnerable to stagnation, cycling, disengagement?
- Three different ideas:
 - Population structure
 - Fitness sharing
 - Reduced virulence

Population Structure

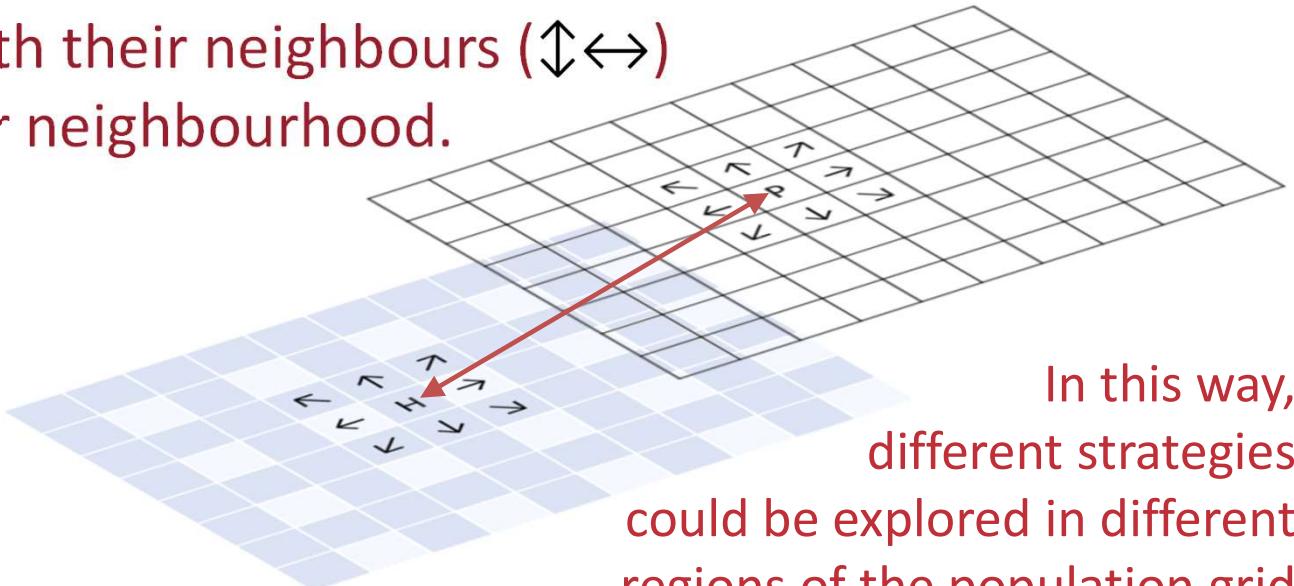
- The natural world is not randomly mixed up like an evolving soup.
- Populations are structured such that not every organism can interact with or compete directly with every other organism
- This is why island ecologies can be so distinctive and interesting
- GAs can do something similar
 - split populations into quasi-isolated sub-populations called ‘demes’
 - spread a population over a space and only allow local interactions



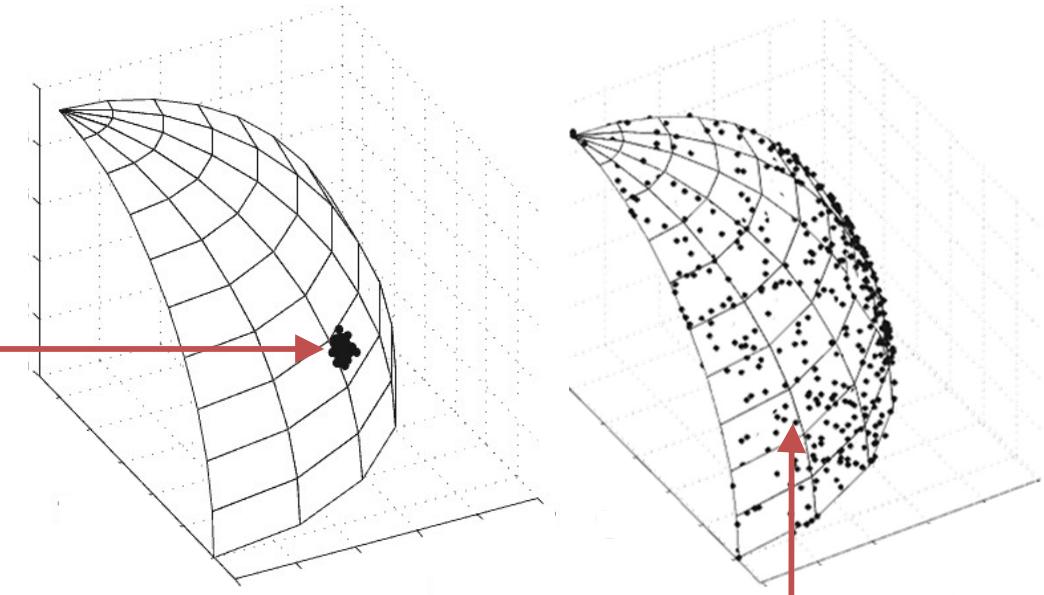
Population Structure

- Hillis employed this idea in his list sorting work:
 - Each host was placed on a grid. Parasites were placed in a different grid.
 - Each host played the parasite at the same location in the other grid (\leftrightarrow)
 - Individuals competed with their neighbours ($\uparrow\downarrow\leftrightarrow$) to leave offspring in their neighbourhood.

H	H	H	H	H	H	H	H
H	H	R	↑	↗	H	H	H
H	H	←	H	⇒	H	H	H
H	H	↖	↓	↘	H	H	H
H	H	H	H	H	H	H	H
H	H	H	H	H	H	H	H



- A more deliberate way of maintaining diversity is to punish solutions that are similar to other solutions in the population:
 - Define a similarity measure (e.g., genetic hamming distance < threshold)
 - If solution i is similar to n others:
 - $f'_i = f_i / \sum_{j=1}^{j=n} \text{sim}(i, j)$
 - Now, instead of converging on one region of this manifold of equally fit solutions...
 - ...fitness sharing pushes solutions to spread out

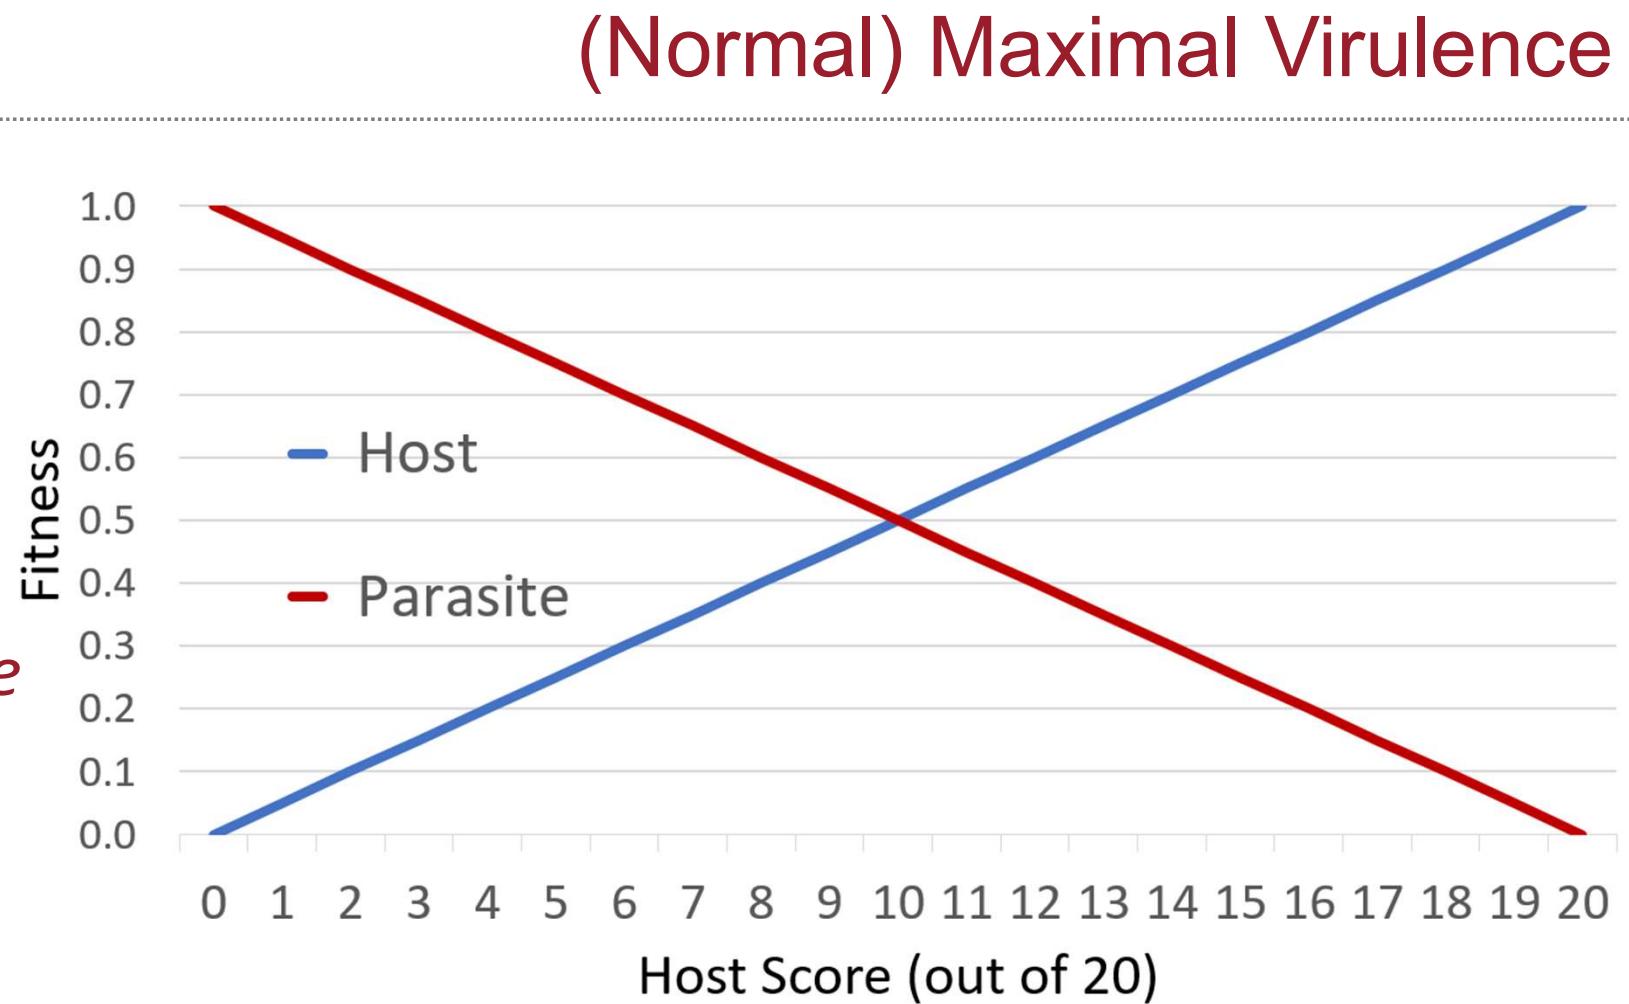


Reduced Virulence

- Natural parasites / viruses often reduce their *virulence* over time
 - ...it's not in their interest to kill their host quickly.
 - Consider the common cold or a tapeworm vs. extremely virulent Ebola
- But parasites in a coevolutionary GA are often rewarded for being *maximally damaging* to the hosts that they are assessed against...
- Our initial random robot goalkeepers won't benefit from facing penalty takers that smack the ball into the top corner every time.
- It would be better to start off with shots that are easy to save, and then get gradually harder...

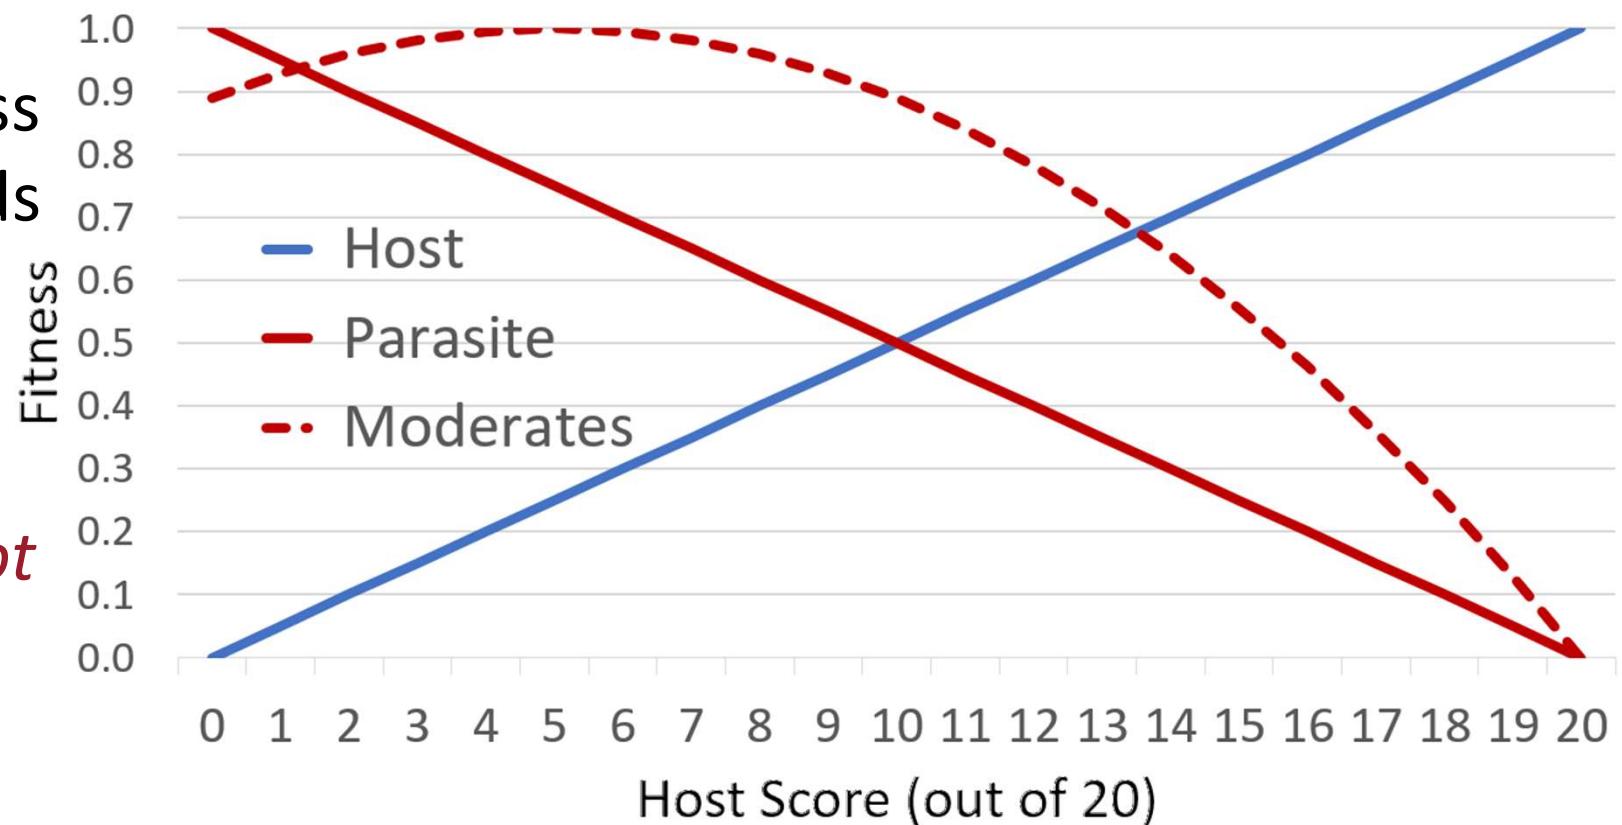
Normally, when a parasite plays against a host:

- If the *Host* scores x ...
- ... the *Parasite* scores $1-x$



Reduced Virulence

A ‘moderate virulence’ fitness function rewards parasites most for defeating hosts *most of the time, but not all of the time...*



- More on these ideas in this weeks' GA Lab
- Ahead of the lab, please take a look at
 - Cartlidge, J. & Bullock, S. (2004). Combating coevolutionary disengagement by reducing parasite virulence. *Evolutionary Computation*, **12**(2), 193-222.
 - This lays out the details of the virulence idea, and introduces a simple game that we will be using to explore it.
 - It's a long paper, so don't feel you need to deep read all of it. ☺
 - Just focus on sections 1-4 (15 pages)

Free Lunches?

- Recall Wolpert & Macready's (1995) No Free Lunch theorem
 - No search alg. can do better than random across *all* problems.
- More recently, in 2005, the same authors claim that *coevolutionary* free lunches are possible...
 - By using a separate (non-random) population to steer its search, a search algorithm may be able to always rule out some bad solutions, and thus out-perform random search...
- Check this stack exchange [post](#) for some intuitions...

Example Questions

- Give an example of coevolution from nature. [1 mark]
 - Name 3 problems that beset coevolutionary GAs. [2 marks]
 - Why did Hillis embed his populations on a grid? [4 marks]
 - What is premature convergence in a GA population. How do coevolutionary GAs hope to avoid it? [4 marks]
 - Explain why a coevolutionary GA might cycle through the same poor solutions over and over again, rather than finding a better general purpose solution. [6 marks]
-

Thank you!

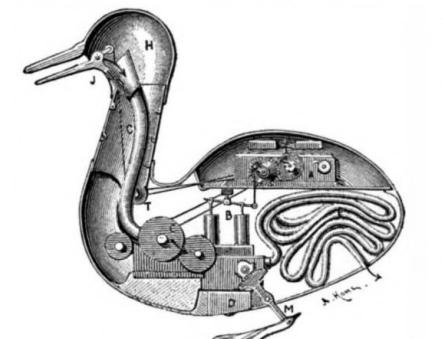
Artificial Life

Seth Bullock

bristol.ac.uk

- Artificial Life is a sister discipline for Artificial Intelligence
- Both take a functionalist approach to a mysterious subject:
 - AI: an artificial system can be *intelligent* if it is organised right
 - Alife: an artificial system can be *alive* if it is organised right
- Like AI, Alife combines a Strong project and a Weak project
 - Strong Alife: creating life inside a computer or in a robot
 - Weak Alife: understanding life by building life-like systems
 - Alife's pithy strapline: "The study of life-as-it-could-be"

- Chris Langton named + founded Artificial Life in 1987
- But Alife has been going on for much longer than AI
 - Frankenstein, Golem, various ‘Automata’ and ‘Living Machines’:
 - Vaucanson’s Duck (1739): quacking, drinking eating, “excreting”
- The biggest names in CS were in fact Alifers...
 - Babbage: Evolutionary Models of Miracles
 - Turing: Models of Biological Morphogenesis
 - Von Neumann: Self-Reproducing Machines

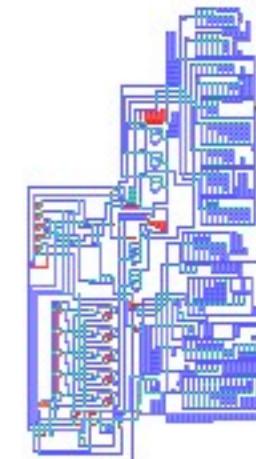


INTERIOR OF VAUCANSON'S AUTOMATIC DUCK.
A, clockwork; B, pump; C, mill for grinding grain; F, intestinal tube;
J, bill; H, head; M, feet.

Self-Reproducing Machines

- How could a machine make a more complex copy of itself?
- This conundrum sounds impossible – paradoxical
- Von Neumann arrived at a two-part solution in 1948:
 1. Information encoding the structure of a machine ...DNA
 2. Machinery that uses the information to make a new machine + a new copy of the information ...cellular machinery
- In doing so, JVM predicted the contents of our cells...
...before the 1952 discovery of the DNA double helix.

- But John Von Neumann went further than theorising and actually designed the first self-replicating machine!
- The first “Cellular Automaton” (CA):
 - A grid of cells, each containing a value
 - All values are updated according to a set of rules applied to each cell grid
 - Careful design of the rules plus the initial state of the grid, and...

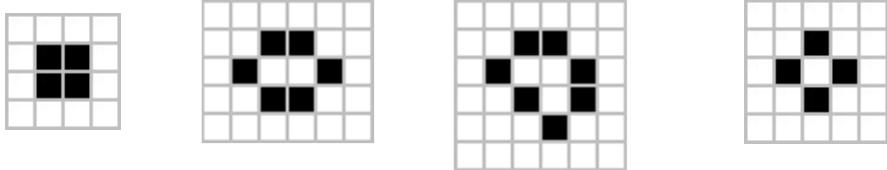


Conway's Game of Life

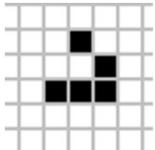
- You may be familiar with the most famous CA: Game of Life
- John Horton Conway (1937-2020) created the *Game of Life* in an effort to make a CA that had interesting behaviour:
 - Proliferating, repeating patterns; universal construction; etc.
- He settled on the following simple rules:
 - Each cell has a binary state – it can be alive or it can be dead
 - A live cell with >3 living neighbours or <2 living neighbours dies
 - A dead cell with three living neighbours becomes alive

Conway's Game of Life

- Some configurations are static:



- Some oscillate on the move:

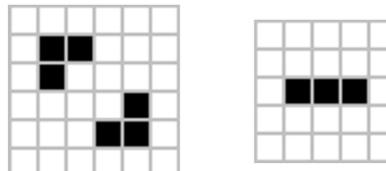


- A “glider”

- Some destroy each other when they collide in the right way...

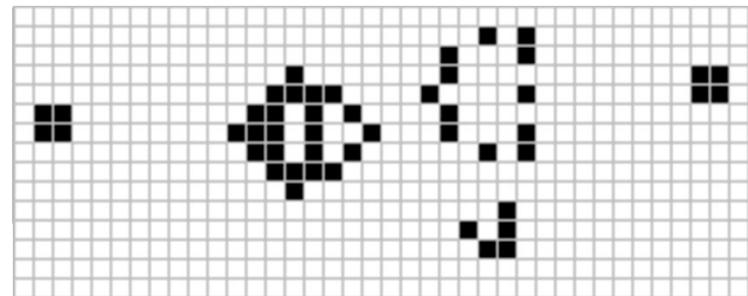
- Amazingly, the Game of Life is provably Turing Complete...

- Some oscillate in place:



- A “beacon” and a “blinker”

- Some *generate* other structures: e.g., a *glider gun*

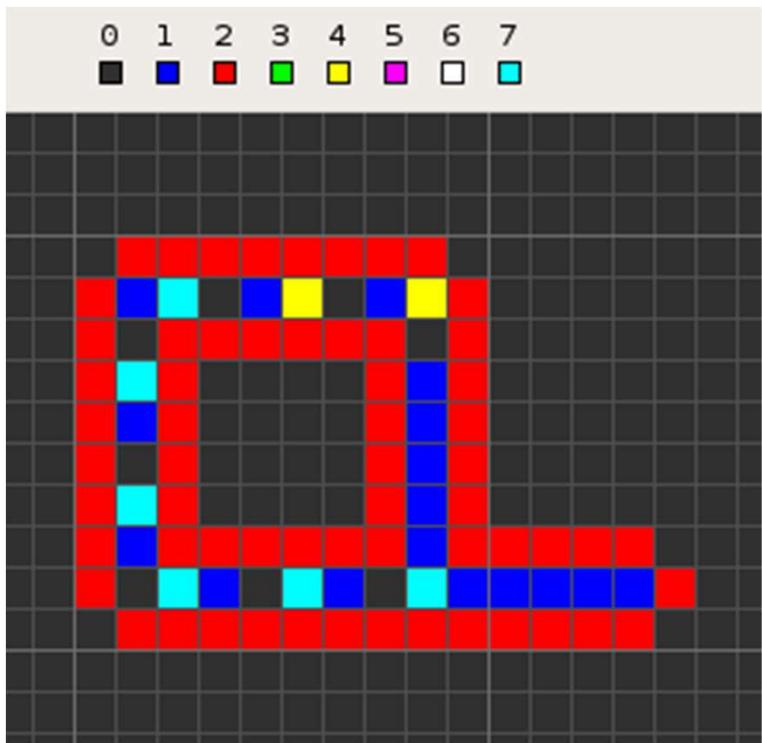


Life in Life
Phillip Bradbury
Roger Pincombe

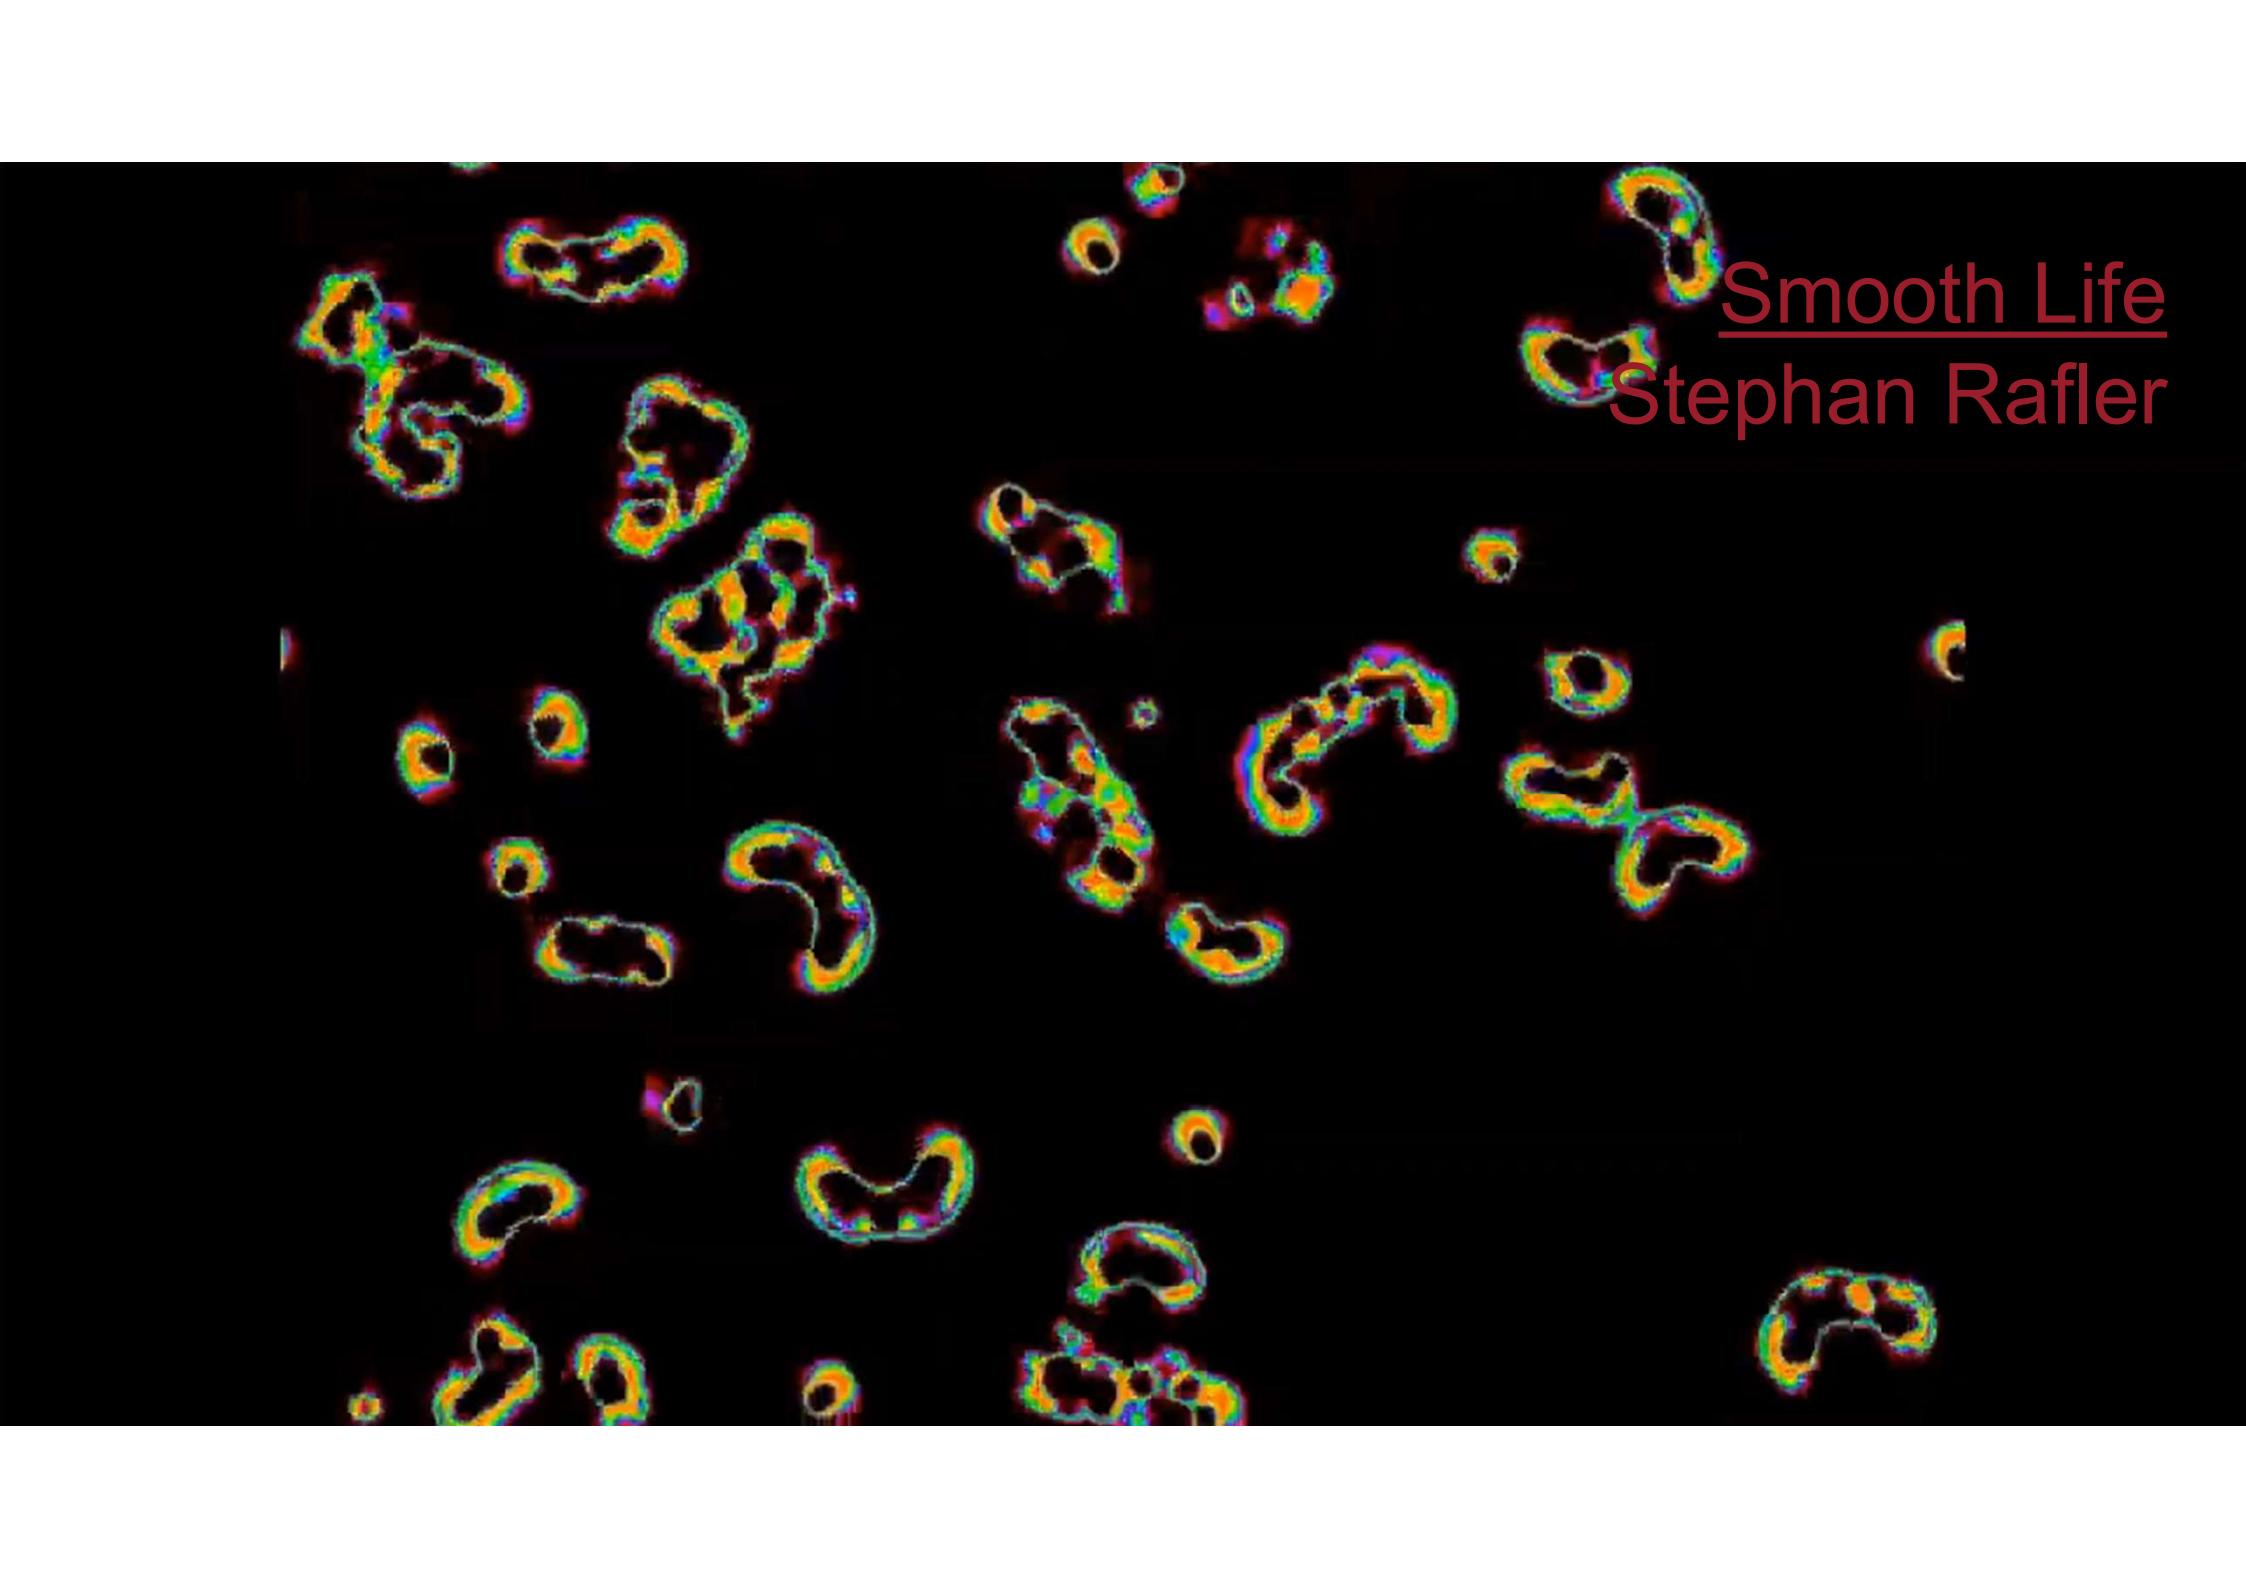


Langton's Loop

- Langton studied CAs + made a more simple self-replicator:



- Some CAs are boring, dead, inert, fixed;
...some just repeat; some are messy chaos;
...and some are ‘interesting’ for a long time
- What is the physics of this complexity?
- Self-replicators + some kind of mutation =
the possibility of *evolving replicators*...
- But how can truly ‘open-ended evolution’
be achieved in artificial systems?



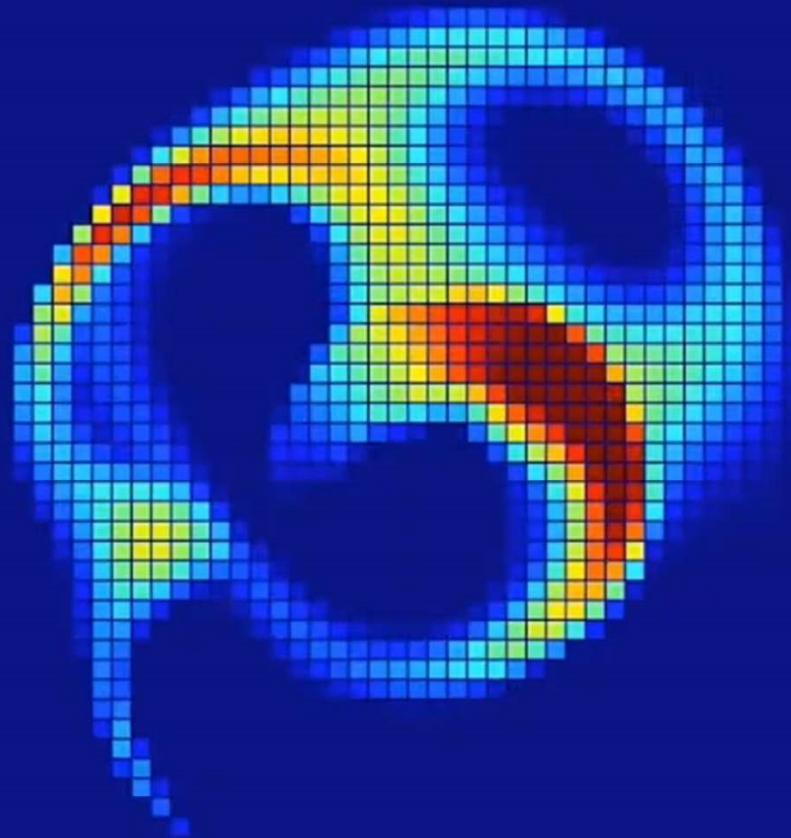
Smooth Life

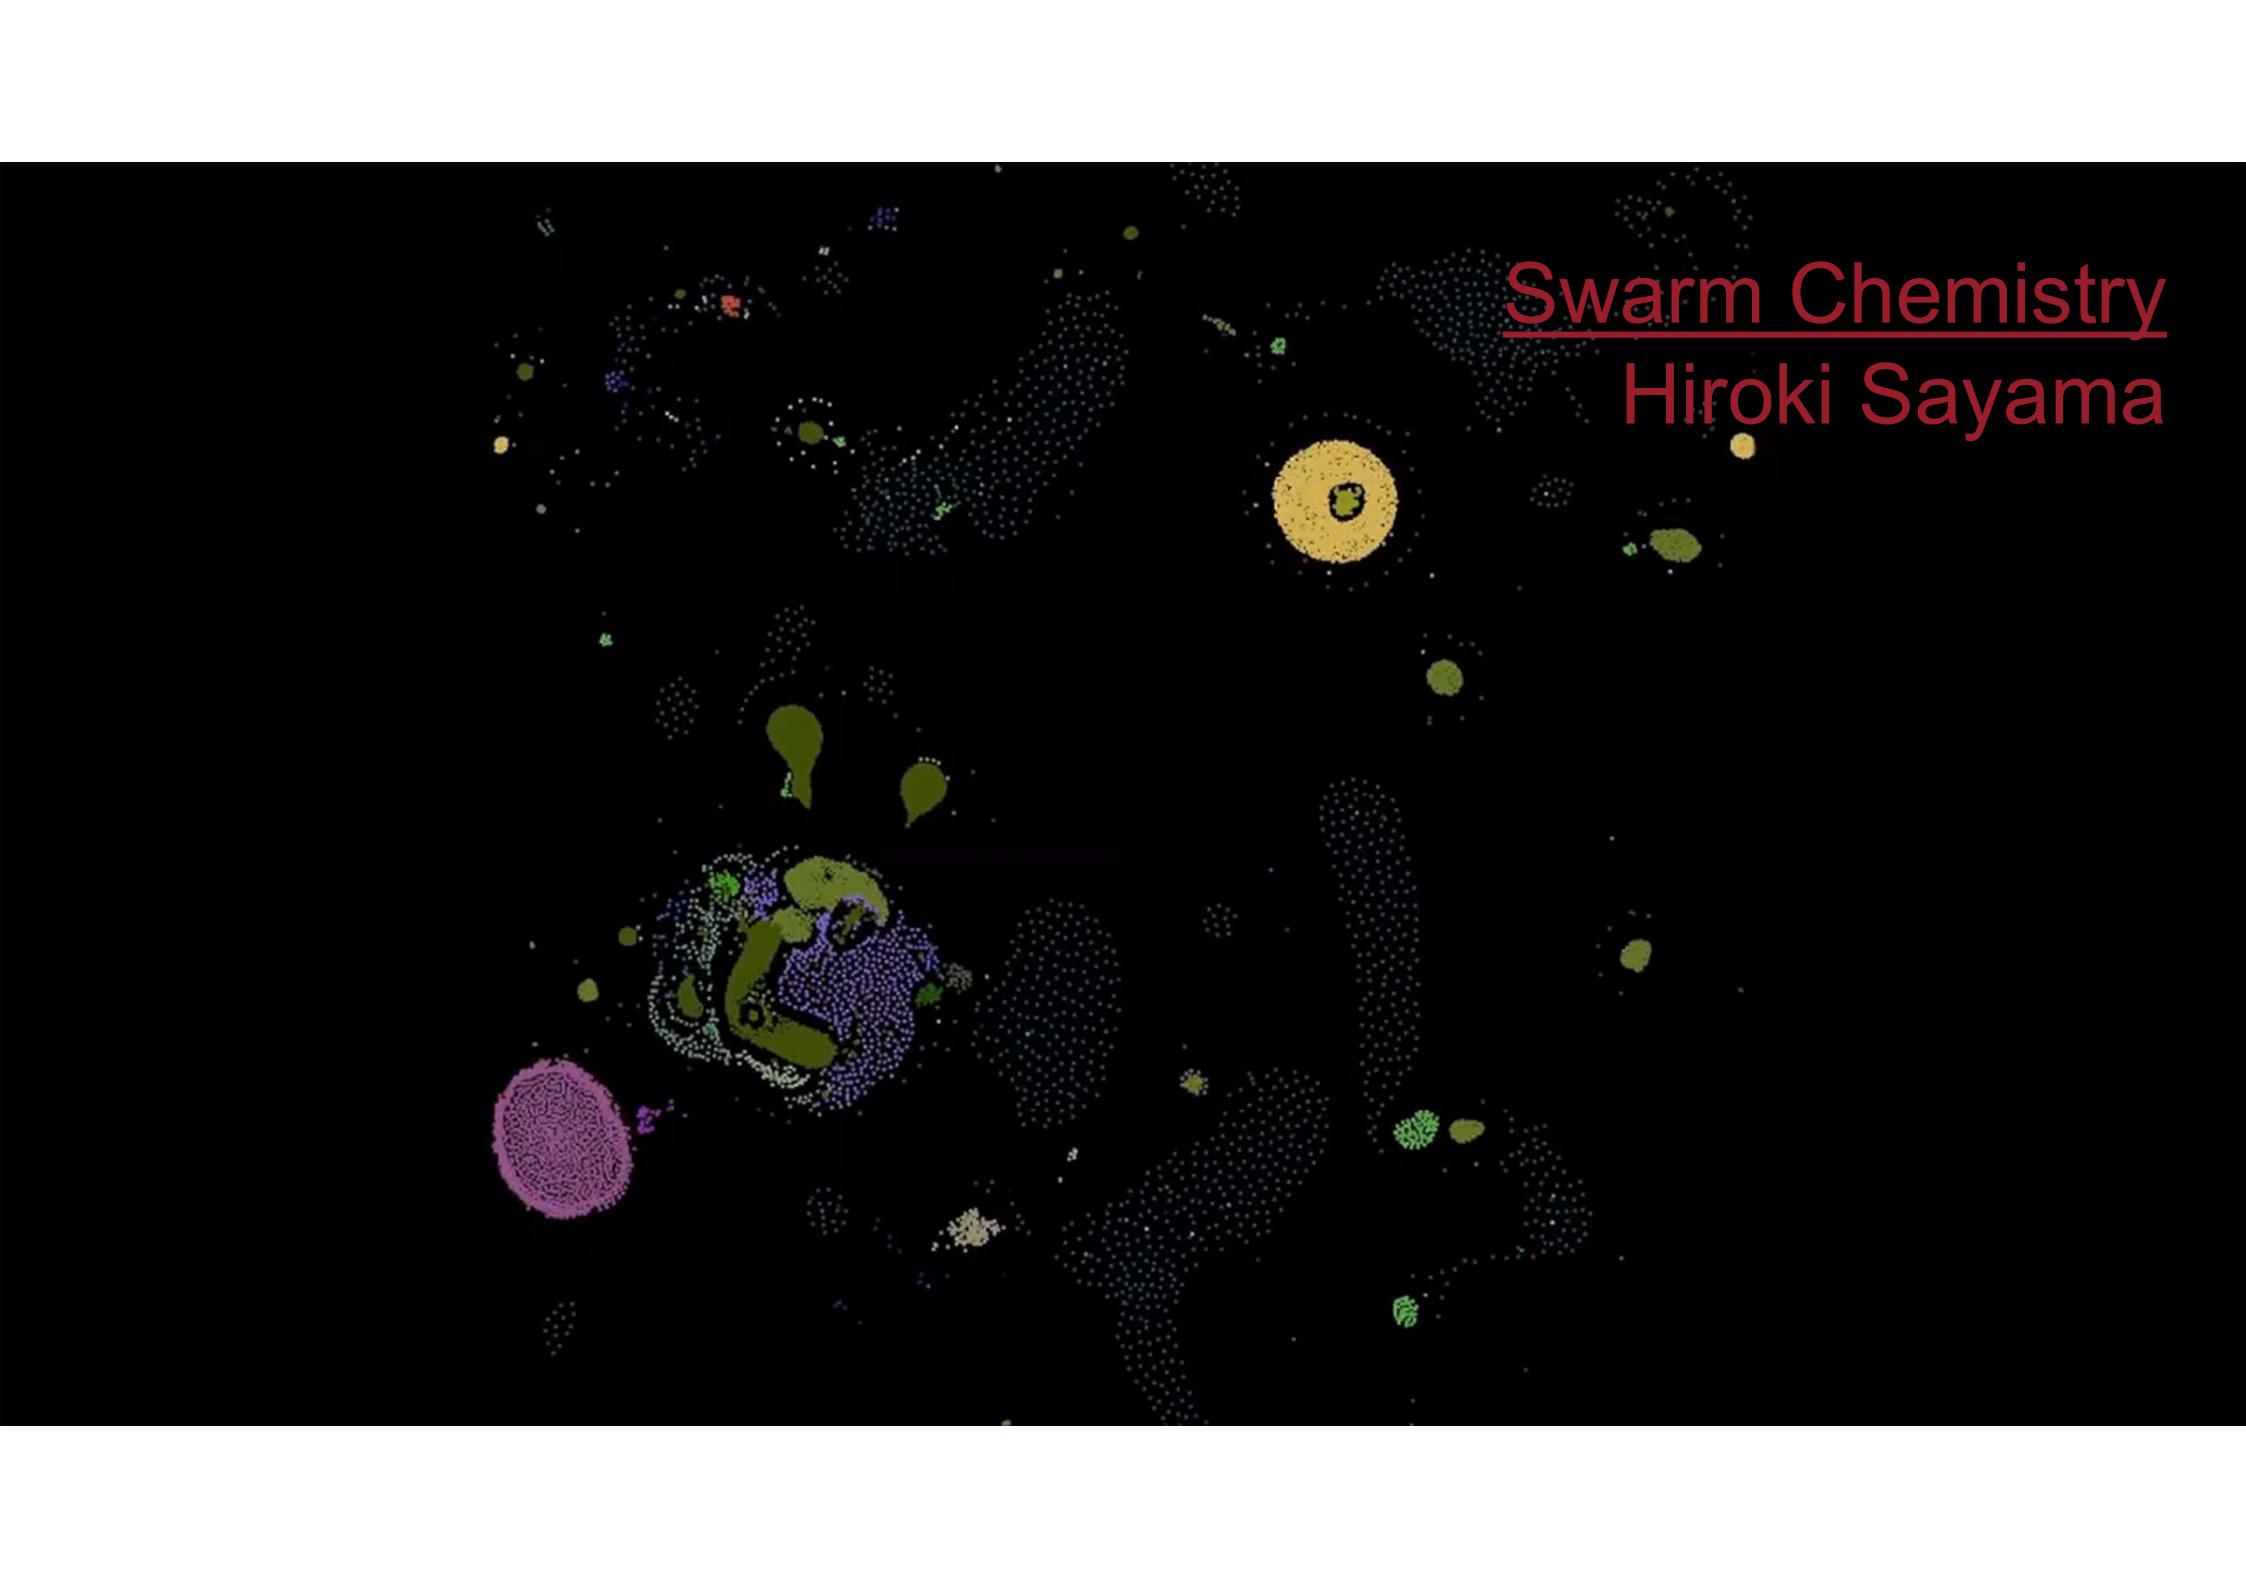
Stephan Rafler



Smooth LifeL
(Tim Hutton)

Lenia
Bert Chan

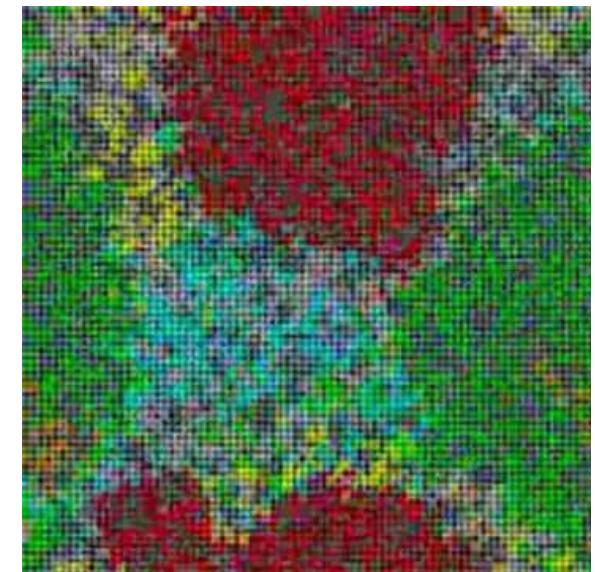
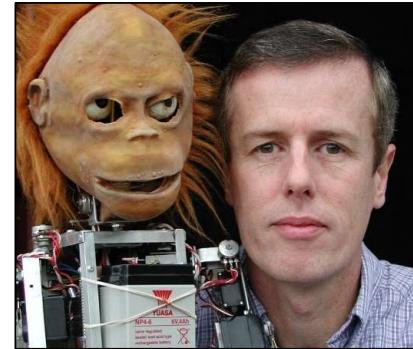




Swarm Chemistry

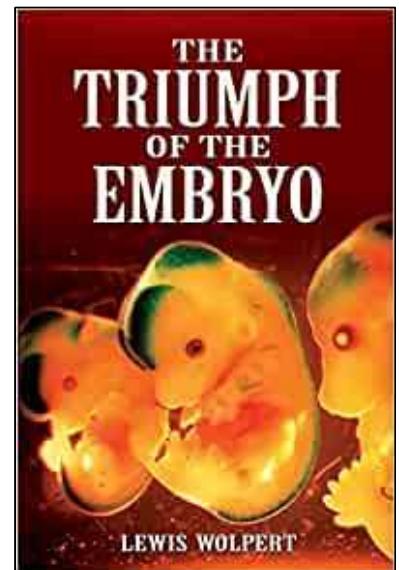
Hiroki Sayama

- Other Alifers explored the evolution of “digital organisms”:
 - Tom Ray’s *Tierra*: from replicators to parasites to hyperparasites
 - Chris Adami’s *Avida*: a kind of digital evolutionary laboratory
 - Alastair Channon’s *Geb*; Yaeger’s *Polyworld*.
 - Steve Grand’s video game: *Creatures*

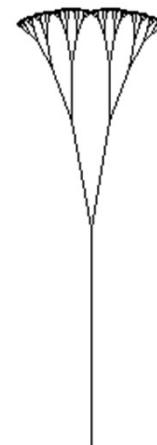


From Evolution to Development...

- Real biological lifeforms *develop* over time: morphogenesis
- Understanding how exactly an egg grows into a person is still one of the most profound scientific challenges
- Since before Turing: *growing* an AI has looked far more feasible than hand-crafting one
 - How can artificial systems grow and develop?
 - How can this development be steered or guided?



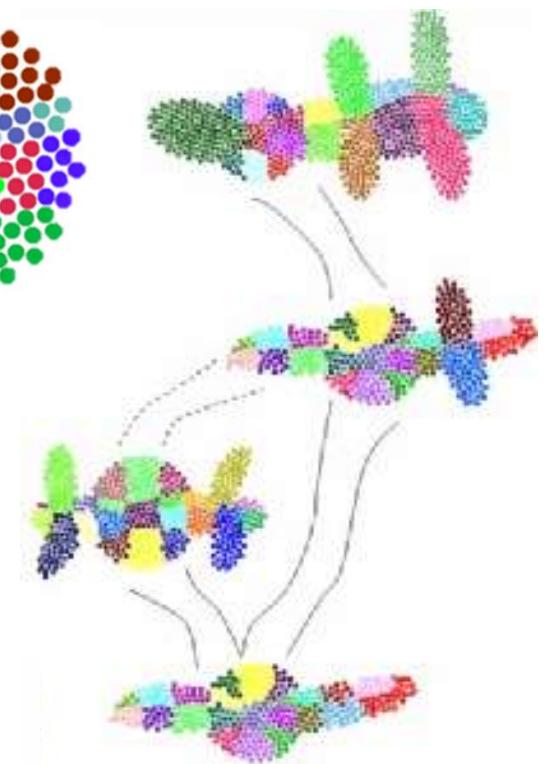
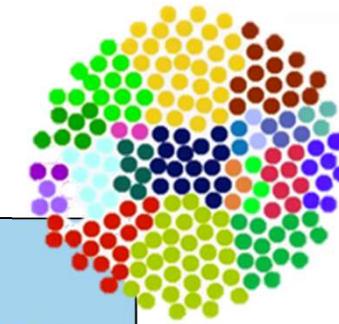
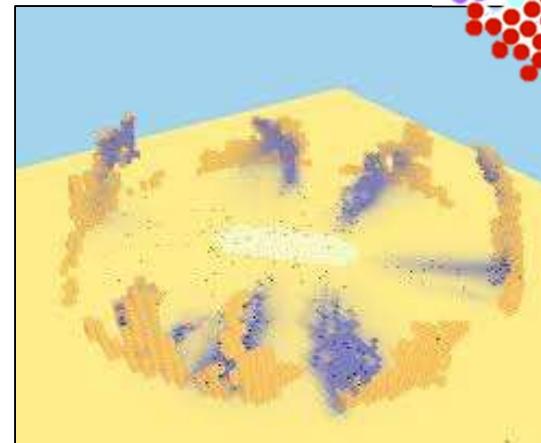
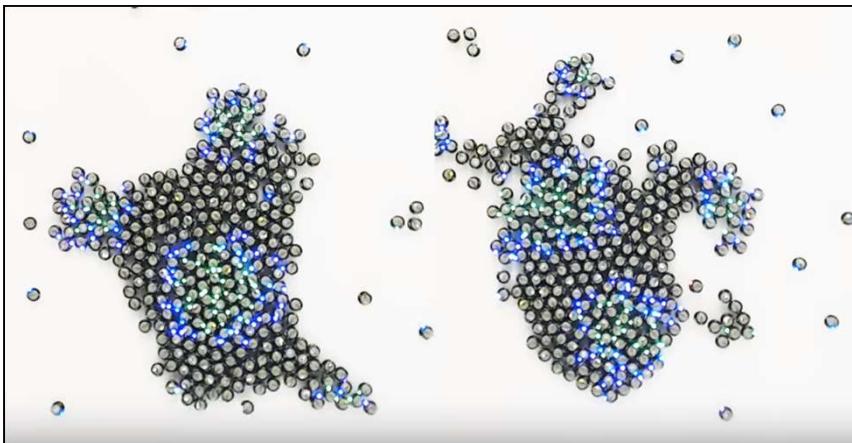
- Lindenmayer Systems, Aristid Lindenmayer (1925-'89)
 - A model of plant development based on grammar rewriting
 - In language: $S \rightarrow NP+VP$; $VP \rightarrow V+NP$; etc., etc.
 - In L-Systems: $I \rightarrow Y$
 - More complex rules:



Przemysław Prusinkiewicz, A. Lindenmayer (1990). *The Algorithmic Beauty of Plants* bristol.ac.uk

Artificial Morphogenesis

- Morphogenesis is still a growing area of research in Alife: 😅
 - “Morphogenetic Engineering”
 - Swarm Construction

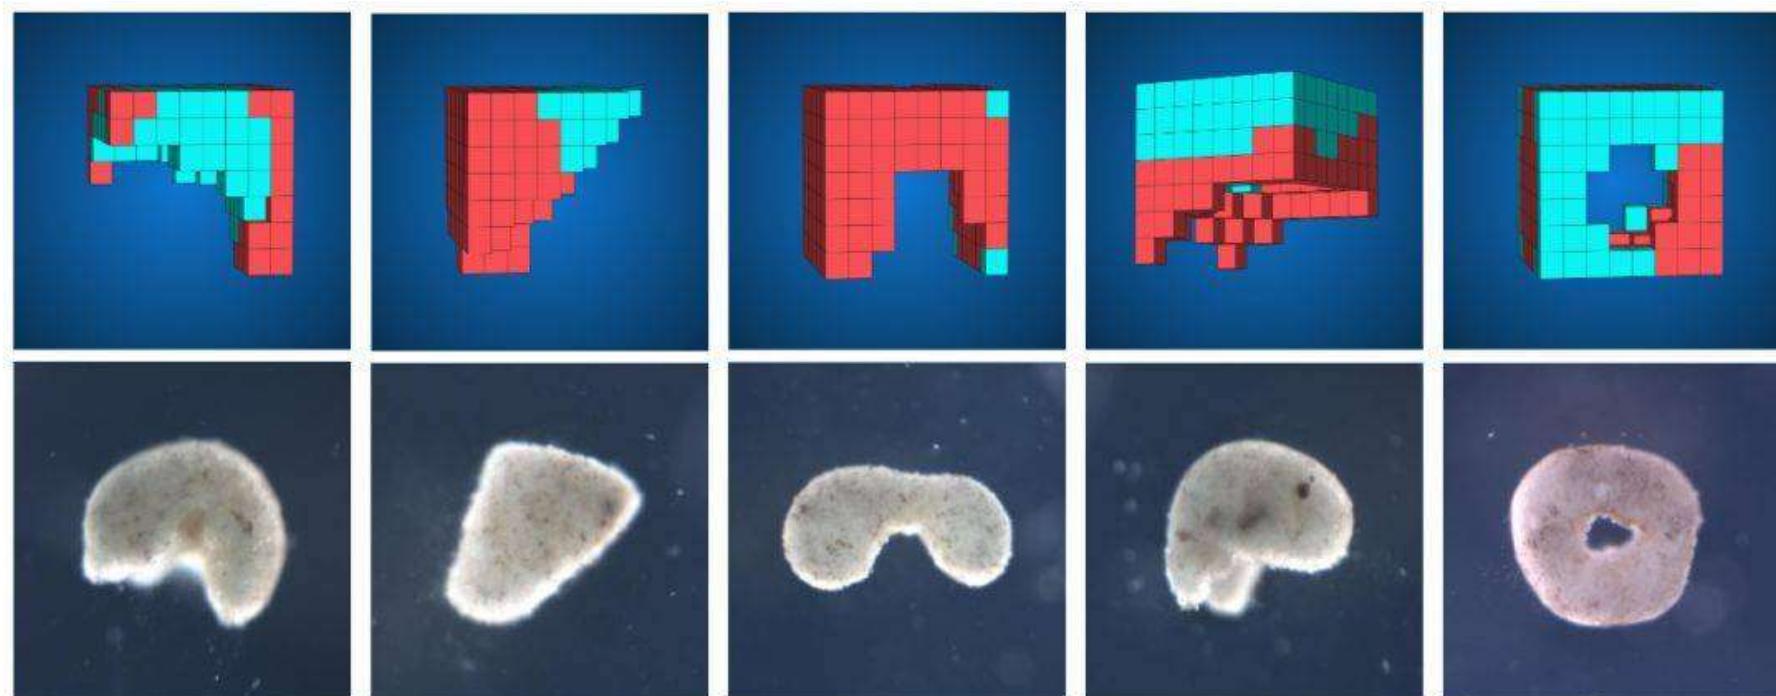


- Evo-Devo Soft Robotics

1



Robotics + soft-body simulation + evo-devo + bio-chemistry...



Josh Bongard (Vermont), Michael Levin (Tufts) and collaborators

bristol.ac.uk

Types of Alife

Alife combines many disciplines ... and considers different life forms:

- Computer Science
 - Robotics
 - Biology
 - Physics
 - Chemistry
 - Philosophy
 - Linguistics
- + more

- Soft Alife: digital life
- Hard Alife: robotic life
- Wet Alife: biochemical life
- Whereas AI has focussed on:
 - Reasoning, Planning, Logic, etc.
- Alife has focussed on:
 - (Co-)Evolution, Development, Learning, Self-Organisation, Emergence, Origins of Life, etc.

Open Problems

In 2000, a group of nine prominent Alife researchers authored a journal paper: “Open problems in Artificial Life”: 14 challenges that are still live + relevant...

- How does life arise from the non-living? (5 problems)
 - Achieve the transition to life in an artificial chemistry.
 - Simulate a unicellular organism over its entire life cycle.
- What are the potentials and limits of living systems? (5 problems)
 - Determine what is inevitable in the open-ended evolution of life.
 - Develop a theory of information processing, flow and generation for evolving systems.
- How is life related to mind, machines, and culture? (4 problems)
 - Demonstrate emergence of intelligence in an artificial living system.
 - Establish ethical principles for artificial life.

Women in Alife!

Pauline
Hogeweg

Janet Wiles

Alex Penn

Jitka Čejková

Katie Bentley

Mizuki Oka

Sabine
Hauert



Charlotte
Hemelrijk



Barbara
Webb



Lana
Sinapayen



Emily
Dolson



Melanie
Mitchell



Hemma
Philamore



Susan
Stepney

Further Reading

- Langton, C. (1989). [Artificial life](#), in *Artificial Life*, Addison-Wesley.
- Lehman, Clune, Misevic, Adami, et al., [52 authors!] (2020). [The surprising creativity of digital evolution: A collection of anecdotes from the evolutionary computation and artificial life research communities](#). *Artificial Life*, 26:2, 274-306.
- International / European Alife Conference proceedings at [alife.org](#)
- The *Artificial Life* journal at [MIT Press](#) | Lana [Sinapayen's Prezi](#) 🔥
- Hiroki Sayama's [PyCX](#) Python package and associated [text book](#)
 - [NetLogo Life](#), [Alife Virtual Seminars](#),  [Alife Papers](#)

Example Questions

- Which cells in the Game of Life grid shown below will survive to the next time step, and which will not? *[2 marks]*
- How did Von Neumann's insight into self-replicating machines anticipate the DNA molecule? *[4 marks]*
- Why might it be more feasible to grow a general purpose AI rather than code one by hand? *[4 marks]*
- Apply two iterations of the following L-system rules to the starting shape shown below. *[3 marks]*

Thank you!